

# Kubernetes Learning Series

## Part 10: Ingress (Practical Implementation)

A Step-by-Step Beginner-Friendly Hands-On Guide  
with Retrospectives, Debugging, and Learnings

By **Sarthak Srivastava**

# Host-based Ingress with NGINX - Beginner-friendly guide

This short guide shows exactly how to route **blue.example.com** and **green.example.com** to **two tiny apps** using an **NGINX Ingress controller** in a **KodeKloud playground Kubernetes cluster**. Follow the steps in the "**Quick steps**" section to get **Hello from Blue!** and **Hello from Green!** working, then read the "**Common mistakes & fixes**" and the compact **troubleshooting** checklist.

---

## What you'll get

- Copy-paste-ready commands to run in order.
  - Beginner-friendly explanations (what and *why*).
  - The common mistakes I made and how to avoid them.
  - A short checklist for fast debugging.
- 

## Prerequisites

- kubectl configured to your cluster.
  - A playground or local Kubernetes cluster (KodeKloud, minikube, kind).
  - Edit access to /etc/hosts on your machine or VM that will run curl.
- 

## Quick steps (copy & run to produce the correct output)

Run these commands in order on a fresh cluster. After step 6 you'll see Hello from Blue! and Hello from Green!

### 1) Deploy the apps (Blue & Green)

**Save and run the following** (creates deployments + services):

```
cat > blue-green.yaml <<'EOF'
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

name: blue-deployment

spec:

replicas: 1

selector:

matchLabels:

app: blue

template:

metadata:

labels:

app: blue

spec:

containers:

- name: blue

image: hashicorp/http-echo

args:

- "-text=Hello from Blue!"

ports:

- containerPort: 5678

---

apiVersion: v1

kind: Service

metadata:

name: blue-service

spec:

selector:

app: blue

ports:

- port: 80

targetPort: 5678

---

apiVersion: apps/v1

kind: Deployment

metadata:

name: green-deployment

spec:

replicas: 1

selector:

matchLabels:

app: green

template:

metadata:

labels:

app: green

spec:

containers:

- name: green

image: hashicorp/http-echo

args:

```
- "-text=Hello from Green!"
```

```
ports:
```

```
- containerPort: 5678
```

```
---
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: green-service
```

```
spec:
```

```
  selector:
```

```
    app: green
```

```
  ports:
```

```
    - port: 80
```

```
      targetPort: 5678
```

```
EOF
```

```
$ kubectl apply -f blue-green.yaml
```

```
$ kubectl get pods,svc
```

```
controlplane ~ → kubectl apply -f blue-green.yaml
```

```
kubectl get pods,svc
```

```
deployment.apps/blue-deployment created
```

```
service/blue-service created
```

```
deployment.apps/green-deployment created
```

```
service/green-service created
```

NAME	READY	STATUS	RESTARTS	AGE
pod/blue-deployment-7895957d9b-lswgq	0/1	ContainerCreating	0	2s
pod/green-deployment-c476dbbd7-rxlwx	0/1	ContainerCreating	0	2s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/blue-service	ClusterIP	172.20.41.141	<none>	80/TCP	2s
service/green-service	ClusterIP	172.20.100.217	<none>	80/TCP	2s

**Quick verification** (confirms the app is healthy):

```
# replace <blue-pod> with the actual pod name
$ kubectl port-forward pod/<blue-pod> 8081:5678 &
$ curl http://127.0.0.1:8081 # -> Hello from Blue!
```

```
# stop the port-forward after test
$ pkill -f "kubectl port-forward pod/<blue-pod>"
```

If the curl shows **Hello from Blue!**, the app and service are working.

---

## 2) Install NGINX Ingress controller

```
$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.13.2/deploy/static/provider/cloud/deploy.yaml
```

```
controlplane ~ ➔ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.13.2/deploy/static/provider/cloud/deploy.yaml
kubectl get pods -n ingress-nginx -w
namespace/ingress-nginx created
serviceaccount/ingress-nginx created
serviceaccount/ingress-nginx-admission created
role.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrole.rbac.authorization.k8s.io/ingress-nginx-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
configmap/ingress-nginx-controller created
service/ingress-nginx-controller created
service/ingress-nginx-controller-admission created
deployment.apps/ingress-nginx-controller created
job.batch/ingress-nginx-admission-create created
job.batch/ingress-nginx-admission-patch created
ingressclass.networking.k8s.io/nginx created
validatingwebhookconfiguration.admissionregistration.k8s.io/ingress-nginx-admission created
```

```
$ kubectl get pods -n ingress-nginx
```

```
controlplane ~ ✗ kubectl get pods -n ingress-nginx
```

NAME	READY	STATUS	RESTARTS	AGE
ingress-nginx-controller-7df9fc45fd-z95dw	1/1	Running	0	4m29s

# wait until controller pod is **1/1 Running**

**Check the controller service:**

```
$ kubectl get svc -n ingress-nginx
```

Note: In playground/non-cloud environments **LoadBalancer** usually shows EXTERNAL-IP <pending>, that's expected. Keep going.

---

### 3) Create the host-based Ingress

**Save and apply:**

```
cat > host-ingress.yaml <<'EOF'
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: host-ingress
spec:
  ingressClassName: nginx
  rules:
  - host: blue.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: blue-service
            port:
              number: 80
  - host: green.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: green-service
            port:
              number: 80
EOF
```

```

controlplane ~ → cat > host-ingress.yaml <<EOF
> apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: host-ingress
spec:
  ingressClassName: nginx
  rules:
  - host: blue.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: blue-service
            port:
              number: 80
  - host: green.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: green-service
            port:
              number: 80
> EOF

```

**\$ kubectl apply -f host-ingress.yaml**

**\$ kubectl get ingress**

```

controlplane ~ → kubectl apply -f host-ingress.yaml
kubectl get ingress
ingress.networking.k8s.io/host-ingress created

```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
host-ingress	nginx	blue.example.com,green.example.com		80	1s

**\$ kubectl describe ingress host-ingress**

```

controlplane ~ → kubectl describe ingress host-ingress
Name:             host-ingress
Labels:           <none>
Namespace:        default
Address:
Ingress Class:    nginx
Default backend:  <default>

```



```
Rules:
  Host                Path  Backends
  ----                -
  blue.example.com    /    blue-service:80 (172.17.1.5:5678)
  green.example.com   /    green-service:80 (172.17.1.6:5678)
Annotations:          <none>
Events:
  Type    Reason    Age    From                Message
  ----    -
  Normal  Sync      14m    nginx-ingress-controller  Scheduled for sync
```

describe should show **blue-service:80** (IP:5678) etc.

#### 4) If EXTERNAL-IP is <pending>, convert to NodePort

In many labs the controller service will be type **LoadBalancer** but **no external IP** is provided. Convert it to **NodePort**:

```
$ kubectl edit svc ingress-nginx-controller -n ingress-nginx
```

```
# change: type: LoadBalancer -> type: NodePort
```

```
controlplane ~ → kubectl edit svc ingress-nginx-controller -n ingress-nginx
service/ingress-nginx-controller edited
```

```
$ kubectl get svc -n ingress-nginx
```

```
# note nodePort for port 80, e.g. 80:31535/TCP -> nodePort = 31535
```

```
controlplane ~ ✗ kubectl get svc -n ingress-nginx
```

NAME	AGE	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
ingress-nginx-controller	30m	NodePort	172.20.184.190	<none>	80:31535/T
ingress-nginx-controller-admission	30m	ClusterIP	172.20.133.140	<none>	443/TCP

#### 5) Find node IP and map hostnames

```
$ kubectl get nodes -o wide
```

```
# copy INTERNAL-IP (e.g. 192.168.239.23)
```

```
controlplane ~ → kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-
controlplane	Ready	control-plane	35m	v1.34.0	192.168.239.23	<none>	Ubu
ntu 22.04.5 LTS		5.15.0-1083-gcp		containerd://1.6.26			
node01	Ready	<none>	33m	v1.34.0	192.168.10.236	<none>	Ubu
ntu 22.04.5 LTS		5.15.0-1083-gcp		containerd://1.6.26			

Edit **/etc/hosts** on the machine where you'll run **curl** and add:

```
$ sudo vim /etc/hosts
```

```
controlplane ~ → sudo vim /etc/hosts
```

```
192.168.239.23 blue.example.com
192.168.239.23 green.example.com
```

(Replace 192.168.239.23 with your node IP.)

---

## 6) Final test (this must return the Hello messages)

# use the **nodePort** noted earlier (example 31535)

```
$ curl http://blue.example.com:31535
```

```
$ curl http://green.example.com:31535
```

# -> Hello from Blue!

# -> Hello from Green!

```
controlplane ~ → curl http://blue.example.com:31535
curl http://green.example.com:31535
Hello from Blue!
Hello from Green!
```

# alternative using Host header to node IP:

```
# curl -H "Host: blue.example.com" http://192.168.239.23:31535
```

```
controlplane ~ → curl -H "Host: blue.example.com" http://192.168.239.23:31535
Hello from Blue!
```

If these print the **Hello messages**, **congratulations** , host-based Ingress is working.

---

## II. Common mistakes — what they look like, why they happen, and the exact fix

For each item: **What you see** → **Why it's happening** → **How to fix it (commands / steps)** → **Tip to avoid it**

---

### 1) LoadBalancer shows EXTERNAL-IP: <pending>

**What you see**

```
kubectkl get svc -n ingress-nginx
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
ingress-nginx-controller	LoadBalancer	172.20.x.x	<pending>	80:xxxxxx/TCP

**Why it happens**

- Cloud providers assign an external IP to **LoadBalancer** services.
- Playground / local / on-prem clusters (minikube, kind, KodeKloud) usually don't have a cloud load balancer available, so the IP stays **<pending>**.

### Fix (exact steps)

1. Change the **controller service** to **NodePort**:

```
$ kubectl edit svc ingress-nginx-controller -n ingress-nginx
```

```
# in the editor: change `type: LoadBalancer` -> `type: NodePort`, save
```

2. **Verify**:

```
$ kubectl get svc -n ingress-nginx
```

```
# note nodePort shown as 80:31535/TCP — nodePort is 31535
```

### Tip to avoid

- If you're on a local/playground cluster, plan to use **NodePort** (or minikube/kind-specific add-ons) instead of expecting **LoadBalancer**.
- 

## 2) Hostnames mapped to 127.0.0.1 instead of the node IP

### What you see

- `curl http://blue.example.com:31535` returns nothing or hits the wrong service.
- Your `/etc/hosts` has:

```
127.0.0.1 blue.example.com
```

### Why it happens

- **127.0.0.1** is **localhost** points to your **current machine's loopback interface**.
- The ingress **NodePort** listens on the **node's network IP** (e.g., 192.168.239.23), not necessarily 127.0.0.1.

### Fix (exact steps)

1. Find **node IP**:

```
$ kubectl get nodes -o wide
```

```
# Use INTERNAL-IP, e.g. 192.168.239.23
```

2. Edit `/etc/hosts` and add:

```
192.168.239.23 blue.example.com
```

```
192.168.239.23 green.example.com
```

3. **Test**:

```
$ curl http://blue.example.com:31535
```

# should return **Hello from Blue!**

#### Tip to avoid

- Always map your **test hostnames** to the **node IP** that exposes the **NodePort**
- 

### 3) Not sending the Host: header (or testing incorrectly with port-forward)

#### What you see

- `curl http://127.0.0.1:8080/` returns a default page.
- **Ingress routing** does not send your request to the **right backend**.

#### Why it happens

- **Ingress matches** rules by the **HTTP Host: header**. If the header doesn't match **blue.example.com**, the controller won't route to **blue-service**.
- **Port-forwarding to a service** vs. pod can sometimes **return the controller's own admin/default page** or **hit HTTPS**, causing **confusing** output.

#### Fix (exact steps)

1. Either set **host mapping** or use the domain:

# with `/etc/hosts` mapping in place:

```
$ curl http://blue.example.com:31535
```

Or include the **Host header** explicitly in the curl request:

```
$ curl -H "Host: blue.example.com" http://127.0.0.1:8080
```

2. For **port-forward testing**, prefer forwarding to the **pod**:

```
$ kubectl port-forward -n ingress-nginx pod/<controller-pod> 8080:80
```

```
$ curl -H "Host: blue.example.com" http://127.0.0.1:8080
```

#### Tip to avoid

- Always think “what Host header will NGINX see?” If you're not sure, pass it explicitly with `-H`.
- 

### 4) Service selector / port mismatch → “no active Endpoint”

#### What you see

Controller log or `kubectl describe` shows:

Service "default/blue-service" does not have any active Endpoint

Or `kubectl get endpoints` shows empty for that service.

### Why it happens

- Service **selector labels** don't match **Deployment labels**, or
- Service **targetPort** doesn't match container **containerPort**.

### Fix (exact steps)

#### 1. Inspect service and deployment:

```
$ kubectl get svc blue-service -o yaml
```

```
$ kubectl get deploy blue-deployment -o yaml
```

2. Ensure **spec.selector** in the **Service equals** the **pod labels** (in Deployment → template → metadata → labels).
3. Ensure service has **ports**: - **port**: 80, **targetPort**: 5678 if the container listens on **5678**.
4. Verify **endpoints**:

```
$ kubectl get endpoints blue-service
```

```
# should show podIP:5678
```

### Tip to avoid

- Always name labels deliberately (e.g., **app: blue**) and keep service selectors consistent.

---

## Learnings - simple principles to remember

1. **Ingress is a traffic director, not a server.**  
It routes external HTTP(S) into internal services based on **Host** and **Path** rules.
2. **Host-based routing depends on the HTTP Host: header.**  
If that header doesn't match an **Ingress rule**, routing won't happen.
3. **Service types depend on environment.**  
**LoadBalancer** works on **cloud providers** (they give an IP). Local labs usually need **NodePort**.
4. **NodePort = nodeIP + port.**  
You must reach your **node's IP** at the **nodePort** to hit the **controller** from **outside** the cluster.
5. **Start small: pod → service → ingress.**  
Always confirm the pod works first, then the service, then ingress. This keeps debugging simple.

---

## Troubleshooting - short, step-by-step flow you can follow

Use this exact ordered flow. After each step, stop and confirm the expected output before moving on.

### 1. Pod & service health

```
$ kubectl get pods
```

```
$ kubectl get svc blue-service green-service
```

```
$ kubectl get endpoints blue-service
```

**Expected:** pods **Running**, services **present**, endpoints list podIP:**5678**.

## 2. Direct pod test

```
$ kubectl port-forward pod/<blue-pod> 8081:5678 & # background
```

```
curl http://127.0.0.1:8081
```

```
# -> Hello from Blue!
```

If this fails, fix the Deployment/Container first.

## 3. Ingress resource check

```
$ kubectl describe ingress host-ingress
```

**Expected:** rules show **blue.example.com** → **blue-service:80** (podIP:**5678**).

## 4. Controller alive?

```
$ kubectl get pods -n ingress-nginx
```

```
$ kubectl logs -n ingress-nginx <controller-pod> | tail -n 50
```

Look for:

- Found valid **IngressClass**
- successfully validated configuration, accepting
- Backend **successfully reloaded**

If the controller logs mention no active Endpoint, go back to step 1.

## 5. Is controller reachable externally?

```
$ kubectl get svc -n ingress-nginx
```

If **EXTERNAL-IP** is **<pending>** and you're not on cloud: change **LoadBalancer** → **NodePort** (see earlier fix).

## 6. Node IP & /etc/hosts

```
$ kubectl get nodes -o wide # pick INTERNAL-IP
```

```
# add to /etc/hosts:
```

```
# 192.168.x.x blue.example.com
```

## 7. Final test

```
curl http://blue.example.com:<nodePort> # Should be Hello from Blue!
```

```
curl http://green.example.com:<nodePort> # Should be Hello from Green!
```

---

## Retrospective

We set up two tiny web apps and an NGINX Ingress controller to route traffic by hostname. The apps were fine from the start; we verified that by port-forwarding to the pod. The real problems were **environment-related**: the Ingress controller was configured as a **LoadBalancer** (which needs a **cloud provider** to get a **public IP**), our test **hostnames** were pointing at **localhost** instead of the **node IP**, and we sometimes tested without the **Host: header**. Changing the controller to **NodePort**, mapping hostnames to the node's IP in **/etc/hosts**, and testing with the **correct Host header** got everything working.

The **key** lesson: most Ingress “**failures**” are actually **network** or testing mistakes, not incorrect YAML.